

A Hot Bath: Conserving Water Without Losing Uniformity

Summary

Maintaining a warm bath is a coupled problem of heat loss, replenishment and transport. A hot inlet can improve the mean temperature while leaving distant water cool and sending useful heat directly to the overflow. We therefore minimize added water subject to a lower limit in every cell and upper-temperature and spread limits outside a 0.15 m inlet zone, rather than optimizing an average alone.

We prove a coast-then-hold policy optimal for a well-mixed bath, then use a three-dimensional finite-volume thermal network with coefficients derived from textbook correlations and tied to a published immersion study where one exists (Section 2). The best constant rate is compared with an optimized piecewise-constant schedule; neither is proved optimal among all controls.

In a 164.25 L water-volume scenario lasting 30 minutes, with an initial temperature of 40°C and a 39°C lower limit, the best constant-rate policy adds **24.14 L** at 0.805 L/min from the start, and a 6-segment schedule with a 0.1°C design margin needs **21.48 L**, 11% less. The well-mixed optimum is 16.01 L, while an independent energy argument gives a 15.41 L lower bound for the spatial problem. A finer mesh changes the constant-rate result by 0.11%, and the selected scheduled policy passes independent continuous-time bounds on all three tested meshes.

Across 64 Sobol draws over the stated parameter ranges, 41 yield an accepted constant-rate candidate, needing 12–31 L (5th–95th percentile). Weak mixing often defeats this finite search. Motion improves transport, but the associated extra surface loss can offset the saving. The assumed foam layer reduces constant-flow demand to 10.4 L; this is a scenario, not a measured additive effect.

At equal volume the deep, narrow tub needs 17.48 L; body size, position and temperature further change the response. Mixing-specific six-stage candidates use 27.00 L at $D = 0.0007 \text{ m}^2/\text{s}$ and 20.30 L at $0.0013 \text{ m}^2/\text{s}$, with three-grid continuous checks. The weaker-mixing candidate starts immediately, showing why baseline waiting is not transferable. A synthetic two-probe calibration leaves 178 compatible models; repairing four sampled policy failures needs 23.48 L commanded, plus a separate 6 L calibration pulse. All 534 model-grid envelopes pass, conditionally on the finite set. Independent integration, analytic limits and energy accounting support 17 baseline checks. Limited body-capacity and flow-path alternatives preserve sampled feasibility, not unknown optima. The recommendations remain conditional on calibrated heat loss and mixing; AI assistance is disclosed.

Keywords: thermal network; energy balance; lower bound; optimal control; sensitivity analysis

Contents

1	Define the decision before optimizing	3
2	Physical assumptions and scenario inputs	4
2.1	Where the coefficients come from.	5
3	A transparent well-mixed benchmark	6
4	What can be proved about water use?	7
5	A conservative model in three dimensions.	8
5.1	Spatial evidence: the mean is not the whole bath	9
6	Inlet transport, solver and strategy search	10
7	The spatial result and the mixing penalty	10
7.1	A time-varying schedule uses less water.	12
7.2	What does staying close to the start temperature cost?.	13
7.3	Can a user follow the schedule?	14
7.4	Measurement ambiguity changes the decision	15
8	Separate geometry, size and body effects.	16
9	Motion and a bubble layer can change the policy.	17
10	Validation: independent evidence and remaining error	18
11	Resolution, uncertainty and transfer to a real bath	19
11.1	Do the policies survive a different body or flow path?.	20
12	Conclusions and a policy with clear scope.	21
13	A warmer bath, with less replacement water.	22
14	References and reproducible algorithm	23
	Appendix A. Scenario definitions.	23
	Report on Use of AI	25

1 Define the decision before optimizing

The task is to preserve both warmth and spatial uniformity in an overflowing, unheated tub, and to examine geometry, the bather and motion, and a bubble-bath layer [1]. The report separates physical requirements from preference assumptions. There is no supplied temperature record or measured heat-transfer coefficient to fit.

The decision variables are an inlet flow rate and the time at which a constant trickle begins. The tub is already full: added water displaces an equal volume through the overflow. The horizon is 1,800 s. Our baseline requires every cell average to stay at or above 39°C, and requires the cell averages outside a 0.15 m jet-mixing zone around the inlet to stay at or below 41°C and within an instantaneous spread of 1.5°C. These choices operationalize comfort; they are not medical limits or numbers specified by the problem, and the jet zone is an assumption justified in Section 2.1.

$$J = 1000 \int_0^{t_f} q(t) dt \quad (1)$$

$$T_i(t) \geq T_{\min} \quad (\forall i), \quad T_i(t) \leq T_{\max}, \quad \max_{i \in \Omega} T_i(t) - \min_{i \in \Omega} T_i(t) \leq \Delta \quad (i \in \Omega) \quad (2)$$

Here Ω is the set of cells outside the inlet jet zone, q is in m³/s and J is in litres. We prioritize the least water within the constraints, rather than assigning arbitrary weights to unlike units. Tightening the temperature tolerance is a separate scenario. Zero flow is admitted: if the initial stored heat suffices, using no added water is globally water-minimal.

The coverage map in Table 1 links each requested outcome to the argument or result that answers it.

Table 1. Where each requirement is answered

Requirement	Where answered
Temperature in space and time	Sections 3, 5–7; Figures 1–2
Water-use strategy and its scope	Sections 4, 6–7 and 12
Tub/body geometry, size and temperature	Section 8
Motion and bubble-bath additive	Section 9
Validation and sensitivity	Sections 8–11
One-page non-technical explanation	Section 13

Spatial temperature means a control-volume average. It does not bound the unresolved temperature of a faucet jet or a skin-contact film. This distinction determines which practical conclusions the simulation can support.

Modeling options. A lumped model (Newton cooling) keeps one temperature and cannot see a cold corner. A flow solver resolves velocity but needs a turbulence closure and a faucet jet that nothing available here can constrain. A finite-volume thermal network sits between them: it keeps conservation exact, resolves where heat is lost and delivered, and makes mixing an explicit, testable

parameter [5]. We use the lumped model as a proved benchmark (Sections 3–4) and the network for the decision (Sections 5–11).

The comparison in Table 2 motivates the modeling choice: a thermal network resolves spatial differences without requiring the unobserved velocity field of a flow solver.

Table 2. Modeling options compared

Option	Resolves	Evidence it needs	Role here
Lumped (Newton)	Mean temperature	Two loss coefficients	Proved benchmark
Thermal network	Cells, heat paths, overflow	Loss, body and mixing coefficients	Decision model
Flow solver	Velocity and buoyancy	Turbulence closure, jet data	Not used

2 Physical assumptions and scenario inputs

Each assumption below is paired with its reason and, where it matters, the scenario that tests it. Coefficients cannot be identified from the problem statement alone, so they come from the anchors in Section 2.1.

The baseline in Table 3 separates geometry and comfort assumptions from the coefficients whose physical anchors are examined next.

Table 3. Baseline inputs

Quantity	Baseline / units	Status
Tub $L \times W \times H$	$1.50 \times 0.65 \times 0.23$ m	Assumed geometry
Displaced body volume / area	$0.060 \text{ m}^3 / 1.15 \text{ m}^2$	Assumed
Water density / heat capacity	$1000 \text{ kg/m}^3 / 4180 \text{ J/(kg K)}$	Rounded constants
Air / skin / inlet temperature	$22 / 34 / 50^\circ\text{C}$	Fixed reservoirs
Surface / wall / body coefficient	$25 / 6.5 / 25 \text{ W/(m}^2 \text{ K)}$	Correlation-based (next page)
Mixing diffusivity D	$0.001 \text{ m}^2/\text{s}$	Uncalibrated closure; range tested
Initial / lower / upper limit	$40 / 39 / 41^\circ\text{C}$	Preference scenario
Time / allowed spread	$1800 \text{ s} / 1.5^\circ\text{C}$	Preference scenario

A1. Constant density and heat capacity. Water properties change by under 1% between 39 and 41°C , and the constants are rounded values, not an evaluated IAPWS table [3]. *Reason:* the resulting capacity error is small next to the loss coefficients.

A2. Heat exchange proportional to temperature difference. Surface, shell and body exchange follow Newton’s law with effective coefficients [2]. *Reason:* the temperature span is narrow, so linearization is acceptable. The surface coefficient aggregates convection, radiation and

evaporation, so evaporation is not added twice; evaporative volume change is neglected. The evaporative part is linearized about the room temperature; an independent review estimated that this understates its sensitivity near 40°C by several percent (about 7% on the total), which is small next to the tested range.

A3. Skin held at a fixed temperature. *Reason:* thermoregulation is outside the problem; skin temperature is varied (32 and 36°C) instead of resolved.

A4. Inlet-to-overflow surface stream. The hot water travels along a prescribed surface path to the overflow, so short circuit is possible. The stream runs along the row of cells nearest the centre line (half a cell off it on meshes with an even number of rows). *Reason:* a different inlet geometry would need its own transport network and evidence.

A5. Mixing as one effective diffusivity D . Motion enters through D , with a separate surface-loss scenario. *Reason:* molecular diffusion alone would not represent circulation; D is anchored only by an order-of-magnitude scaling (Section 2.1) and is tested over a range; vertical mixing is tested separately in Section 9.

2.1 Where the coefficients come from

No temperature record exists to fit, so each coefficient is derived from a textbook relation or tied to a published measurement where one exists, and given a range; where nothing exists (mixing, the exposed fraction, the air and skin temperatures) the value is an assumption and is labelled as one. Surface loss sums natural convection above a hot horizontal surface ($Nu = 0.15 Ra^{1/3}$ for $10^7 < Ra < 10^{11}$, length A/P [5]; here $Ra = 1.9 \times 10^7$), linearized radiation (emissivity 0.96) and evaporation by the Lewis analogy. For open water at 40°C in 22°C air at 50% humidity the parts are 4.7, 6.1 and 24.6 W/(m² K) (0.66 kg/m² h evaporated), 35.4 in total. A bather covers part of the surface; with an assumed exposed fraction of 0.7 the central value is 24.8, and we use 25 within the range 17–37, which spans exposed fractions of 0.5 to 1.0. Evaporation correlations differ by tens of percent, so this is a scenario range, not a measured one.

The shell coefficient is a series resistance: a 5 mm shell (0.19 W/(m K)) with film coefficients of 300 inside and 8 W/(m² K) outside gives 6.5, and 4.7–8.4 for other thicknesses. These are typical engineering values, not measurements of a particular tub. The body coefficient is anchored by a measurement: Menzies et al. report a rectal-temperature rise of $0.9 \pm 0.3^\circ\text{C}$ after 30 minutes in 40°C water to the shoulders [6]. Taking the mean body mass of the study participants, 73 kg, $c = 3470 \text{ J/(kg K)}$ and a mean body rise of 1–2°C, the average heat uptake is 141–281 W, equal to 20–41 W/(m² K) with skin held at 34°C. We use 25, because a fixed skin node overstates the driving difference later in the bath. The mean body rise and the driving difference are assumptions, so this is an order-of-magnitude anchor, not a calibration. Holding the skin at a fixed temperature also ignores the heat capacity of the body, about 37% of the water's here (73 kg $\times$ 3470 J/(kg K) against 687 kJ/K): uptake is anchored only on its 30-minute average. Its decay is omitted in the baseline and tested with a finite effective reservoir in Section 11.1.

Mixing has only a scaling anchor. A mixing-length estimate $D \approx 0.1 u' \ell$, with a velocity scale

u' and eddy size ℓ , gives about 1×10^{-4} m²/s for buoyancy-driven flow alone ($u' = 0.02$ m/s, $\ell = 0.05$ m), 1×10^{-3} for gentle movement of the bather (0.1 m/s, 0.1 m) and 6×10^{-3} for vigorous movement (0.3 m/s, 0.2 m). We use 0.001 and test 0.0003–0.003. This is a scaling argument, not a measurement, so D stays a scenario.

The 39–41°C window is a preference; immersion studies use 40–42°C water [6], which is not a safety standard. The inlet cell is a point source: the same policy gives a maximum inlet-cell average of 40.55, 40.76 and 41.00°C on the three meshes of Section 11, because a smaller cell mixes less water with the hot inflow. A limit imposed on that cell would test the mesh, not the bath. The ceiling and the spread limit therefore apply outside a jet-mixing zone of 0.15 m radius around the inlet, on the assumption that a bather does not sit in the jet; the floor applies to every cell, which the energy bound of Section 4 needs. With this definition the constrained temperatures agree across meshes (Section 11).

The relations and ranges in Table 4 show how each heat-loss coefficient is anchored; the selected values define a scenario rather than a fitted bath.

Table 4. Coefficient anchors and the values used

Coefficient	Relation or anchor	Derived value	Used (range tested)
Surface, W/(m ² K)	Convection + radiation + evaporation [5]	35.4 open; 24.8 at 70%	25 (17–37)
Shell, W/(m ² K)	Series resistance, typical shell values	6.5 (4.7–8.4)	6.5 (4.5–8.5)
Body, W/(m ² K)	Uptake implied by core rise [6]	20–41 at fixed skin	25 (12–40)
Mixing D , m ² /s	Mixing-length scaling	1×10^{-4} to 6×10^{-3}	0.001 (0.0003–0.003)

3 A transparent well-mixed benchmark

The notation in Table 5 distinguishes cell capacities, transport and boundary exchange so the energy balance can be checked term by term.

Table 5. Symbols

Symbol	Meaning
C	Heat capacity of the water, J/K
H_a, H_b	Air/shell and body conductance, W/K
T_a, T_b, T_{in}	Room, skin and inlet temperature, °C
T_{min}, T_{max}	Lower and upper limit, °C
q, J	Inlet flow, m ³ /s; added water, L
D	Effective mixing diffusivity, m ² /s

Let C be total water heat capacity, H_a the air/shell conductance and H_b the body conductance. A well-mixed overflow has the same temperature as the bath. Integrating the physical energy balance gives

$$C\dot{T} = H_a(T_a - T) + H_b(T_b - T) + \rho c_p q(T_{in} - T) \quad (3)$$

$$C = \rho c_p (LWH - V_b), \quad H_b = h_b A_b \quad (4)$$

$$H_a = h_s f LW + h_w \{LW + 2H(L + W)\} \quad (5)$$

The foam multiplier f equals one without a layer. Setting $q = 0$, define $H = H_a + H_b$ and $T_e = (H_a T_a + H_b T_b)/H$. The cooling solution and first time to reach the lower limit are

$$T(t) = T_e + (T_0 - T_e)e^{-Ht/C} \quad (6)$$

$$t_c = \frac{C}{H} \log \frac{T_0 - T_e}{T_{\min} - T_e} \quad (7)$$

Baseline conductances are $H_a = 37.14$ W/K and $H_b = 28.75$ W/K, and $C = 686,565$ J/K. Thus $T_e = 27.24^\circ\text{C}$ and $t_c = 14.17$ min.

A heat-loss model must approach an environmental equilibrium rather than zero Celsius. This analytic boundary also supplies a direct check on signs and units. The benchmark describes mixing perfectly; the spatial model will test the cost of departing from that assumption.

4 What can be proved about water use?

Let $\ell(T) = H_a(T - T_a) + H_b(T - T_b)$ be the heat loss of the mixed bath, and let the inlet be hotter than every temperature considered. **Proposition 1.** Coasting to the lower limit and then holding it uses the least water, provided the holding rate is within the faucet bound. The proof has four steps.

Proof. **Step 1, an identity.** Dividing the balance by $T_{in} - T$ and integrating over the horizon gives

$$\int q dt = \frac{C}{\rho c_p} \log \frac{T_{in} - T_0}{T_{in} - T_f} + \int \frac{\ell(T)}{\rho c_p (T_{in} - T)} dt \quad (8)$$

Step 2, monotonicity. The integrand $\ell(T)/(T_{in} - T)$ has derivative $[H_a(T_{in} - T_a) + H_b(T_{in} - T_b)]/(T_{in} - T)^2 > 0$, so it increases with T , and the logarithmic term increases with the final temperature. Water use is therefore an increasing functional of the temperature path.

Step 3, the lowest admissible path. Because $q \geq 0$, comparison with the no-flow solution gives $T(t) \geq T_c(t)$, the cooling curve of Section 3, and feasibility requires $T(t) \geq T_{\min}$. The pointwise lowest admissible path is $\max(T_c(t), T_{\min})$: coast until T_c reaches $T_{\min}$, then hold.

Step 4, attainability. Holding $T = T_{\min}$ means $\dot{T} = 0$, which needs the rate

$$q_{hold} = \frac{\ell(T_{\min})}{\rho c_p (T_{in} - T_{\min})} \quad (9)$$

When this rate is within the bound, the path is feasible and, by Step 2, optimal.

□

For the baseline, the mixed-model optimum adds 16.01 L: wait 14.17 min, then supply 1.012 L/min. Optimality is established for the mixed model; spatial policies are assessed separately.

Proposition 2 (energy lower bound). In the spatial model, suppose every cell satisfies $T_i(t) \geq T_{\min}$. Summing the cell balances gives $\sum_i C_i T_i(t_f) - CT_0 = - \int L dt + \rho c_p \int q(T_{\text{in}} - T_{\text{out}}) dt$, with L the total loss. Each loss term increases with its cell temperature, and H_a and H_b are the sums of the cell conductances, so $L(t) \geq \ell(T_{\min})$. The outlet cell satisfies $T_{\text{out}} \geq T_{\min}$, so $T_{\text{in}} - T_{\text{out}} \leq T_{\text{in}} - T_{\min}$, and the left side is at least $C(T_{\min} - T_0)$. Combining these three facts,

$$J \geq \max \left(0, \frac{1000[\ell(T_{\min})t_f - C(T_0 - T_{\min})]}{\rho c_p (T_{\text{in}} - T_{\min})} \right) \quad (10)$$

The resulting 15.41 L is a genuine conditional lower bound, valid for any control, but it is not tight: the initial water is hotter than the floor and loses more heat, and limited mixing adds further cost. Section 7.1 shows how much of the gap an optimized schedule recovers.

5 A conservative model in three dimensions

The rectangular water envelope is divided into $8 \times 4 \times 3$ cells. Each has a capacity C_i and exchanges heat only across common faces. The cells whose centres lie within 0.15 m of the inlet cell form the jet zone (2, 9 and 22 cells on the $8 \times 4 \times 3$, $12 \times 6 \times 4$ and $16 \times 8 \times 6$ meshes); Ω is the rest. Top cells lose heat to the room; bottom and side faces lose heat through the shell. The bather is represented by a smooth, three-dimensional displacement field and a distributed skin contact term.

$$C_i \dot{T}_i = \sum_j g_{ij}(T_j - T_i) + H_{a,i}(T_a - T_i) + H_{b,i}(T_b - T_i) + S_i \quad (11)$$

$$g_{ij} = \rho c_p D \frac{A_{ij}}{d_{ij}} \min(\phi_i, \phi_j), \quad C_i = \rho c_p V_i \quad (12)$$

Here ϕ is fluid fraction, A_{ij} a face area and d_{ij} the center separation. Symmetric g_{ij} guarantees that internal heat exchange cancels when the equations are summed. D is an effective mixing closure; molecular diffusion alone would not represent motion or buoyant circulation.

$$w_i \propto \exp \left[-\frac{1}{2} \sum_{k=1}^3 \left(\frac{x_{ik} - b_k}{s_k} \right)^2 \right], \quad \sum_i w_i = 1 \quad (13)$$

$$V_i = V_{\text{cell}} - V_b w_i, \quad H_{b,i} = h_b A_b w_i \quad (14)$$

Baseline $\mathbf{b} = (0.55L, W/2, 0.45H)$, with widths $\mathbf{s} = (0.40, 0.16, 0.12)$ m. Fluid fraction is $\phi_i = V_i/V_{\text{cell}}$. The center $\mathbf{b}$ and widths $\mathbf{s}$ set the spatial distribution of body effects. Body volume and contact area remain independent inputs: shape is varied through $\mathbf{s}$ at fixed volume and area. This

is a homogenized immersed-body representation, not an anatomically resolved obstruction. A cell whose occupied fraction reaches 0.95 is rejected rather than assigned a negative capacity.

A fixed envelope volume is preserved by balancing inlet and overflow. In the baseline the exact displaced-water volume is 164.25 L. Shape scenarios preserve envelope volume when isolating aspect-ratio effects. Volume scenarios deliberately change water depth and hence both storage and side-wall area.

5.1 Spatial evidence: the mean is not the whole bath

The layers in Figure 1 locate the remaining spatial temperature differences; a single shared scale makes the top-layer inlet path comparable with the deeper water.




Figure 1. Final temperatures in the three horizontal cell layers. All layers use one color scale; the top layer contains the prescribed inlet-to-overflow stream. Geometry is in metres, and values are cell averages.

The heat map is calculated from the same archived trajectory as Figure 2. It shows where the imposed transport path and environmental/body sinks leave temperature differences. The plots are horizontal slices through a three-dimensional network with exchange between layers, not three independent two-dimensional models.

A temperature range summarizes the spread but cannot show its location. The maps make the physical interpretation inspectable: localized replenishment and distributed losses must be balanced through mixing. Every cell contributes to the comfort test; a high mean cannot compensate for a cold region.

The Gaussian body representation changes both storage and exchange distribution. A map does not prove that this homogenized representation captures anatomy, recirculation or buoyancy. Its role is to reveal the actual implications of the declared model, not to create the appearance of a resolved flow simulation. The refinement table separately checks how cell size affects the result.

6 Inlet transport, solver and strategy search

Inlet and overflow are on opposite ends of a surface stream. Every path edge carries the same q : the first cell receives hot water, interior cells receive upstream water and lose the same volume downstream, and the last cell discharges through the overflow.

$$S_1 = \rho c_p q (T_{in} - T_1), \quad S_i = \rho c_p q (T_{i-1} - T_i) \quad (15)$$

$$\sum_i C_i \dot{T}_i = - \sum_i H_{a,i} (T_i - T_a) - \sum_i H_{b,i} (T_i - T_b) + \rho c_p q (T_{in} - T_{out}) \quad (16)$$

This stream deliberately allows hot-water short circuit. Heat can leave at a locally warm outlet before warming a distant cold region. It is an explicit transport assumption, not an inferred flow field. Section 11.1 changes the internal path while preserving inlet and overflow; alternative inlet placement would still require another network and calibration.

For a fixed flow, the network is affine linear. Augmenting the state by a constant one permits matrix-exponential propagation, avoiding a large forward-Euler step restriction [4]. A delayed-start policy has two exact constant-control segments.

$$\dot{\mathbf{T}} = A(q)\mathbf{T} + \mathbf{b}(q), \quad \mathbf{z}(t + \tau) = e^{\tilde{A}(q)\tau} \mathbf{z}(t) \quad (17)$$

Starts are tested at 0, 120, ..., 1200 seconds. For each, rates from 0 to 3 L/min are bracketed on a 0.2 L/min grid; the first crossing of the lower-temperature requirement is refined by Brent's root method. The search targets $T_{\min} + 0.03^\circ\text{C}$ as numerical reserve, then checks the upper-temperature and spread limits at 5-second output intervals. Starts that already violate the floor are rejected.

This is a reproducible candidate search. Nonmonotone flow responses, untested delays, pulsed inputs, inlet relocation and feedback are not excluded by the calculation. A failed search means no candidate was accepted, not that every possible action is infeasible. The independent checks further examine the chosen trajectory between samples.

7 The spatial result and the mixing penalty

The trajectory in Figure 2 contrasts the volume-weighted mean with the coldest cell and the full spatial range, showing why the mean alone cannot establish comfort.

The results in Table 6 show that the selected constant rate meets the stated temperature limits while accounting for the total replacement water.

Within the constant-rate family, starting the trickle immediately ranks ahead of delayed starts here, even though waiting is optimal in the ideal mixed model. The distant water needs time to receive heat; larger late rates create a warmer inlet region before they solve the cold-region constraint. This is a concrete consequence of resolving space.

Holding a perfectly mixed bath at 40°C throughout would use 36.22 L. This is a stricter reference service, not a like-for-like optimum. Our accepted 1°C cooling allowance uses less water partly




Figure 2. Matrix-exponential cell temperatures under the selected policy. The shaded range covers the coldest cell and the hottest cell outside the inlet jet zone; the mean alone would hide cold and hot locations.

Table 6. Baseline result under the best constant rate

Result	Baseline
Rate / start	0.805 L/min / 0.0 min
Added water	24.14 L
Minimum / maximum	39.030 / 40.318°C
Maximum simultaneous spread	0.967°C
Final volume-weighted mean	39.418°C

because it provides a different service. The comparable mixed model with the same lower limit uses 16.01 L. The gap also reflects our restricted control family and 0.03°C numerical reserve; it cannot be attributed purely to imperfect mixing.

7.1 A time-varying schedule uses less water

Constant-rate policies are a narrow family. We therefore optimize piecewise-constant inlet flow in K equal segments ($K = 3, 6, 12$) by sequential quadratic programming, minimizing added water under the same floor and outside-jet limits (with the 0.03°C floor reserve), sampled every 5 s, from four starting schedules. The problem is non-convex: each result is the best local optimum found, with no global claim.

The comparison in Table 7 separates feasible search results from proved bounds: scheduling saves water relative to the constant rate, but the remaining gap is not an optimality certificate.

Table 7. Constant rate, scheduled flow and the bounds

Policy	Water (L)	Change	Min / max (°C)	Spread (°C)
Best constant rate	24.14	—	39.03 / 40.32	0.97
3 segments, coarse candidate	19.75	-18.2%	39.03 / 40.74	1.50
6 segments, coarse candidate	19.50	-19.2%	39.03 / 40.73	1.50
12 segments, coarse candidate	19.35	-19.9%	39.03 / 40.80	1.50
6 segments, accepted	21.48	-11.0%	39.13 / 40.79	1.42
Perfect-mixing optimum	16.01			
Energy lower bound	15.41			

```

Input: network, limits, K, starting schedules
for each starting schedule x0:
    minimize sum(x)*segment_time over 0 <= x <= 3 L/min
    subject to min_i T_i(t) >= Tmin+reserve,
               max_{i in Omega} T_i(t) <= Tmax,
               spread over Omega <= span (every 5 s)
retain coarse candidates and buffered alternatives;
replay on three meshes with independent integration;
accept only if continuous bounds meet all limits

```

The schedule in Figure 3 concentrates replenishment in the middle of the bath; its temperature trajectory must therefore be checked over the no-flow intervals as well.

The accepted 6-segment policy is off for the first 5.0 minutes and last 10.0 minutes, with a peak flow of 1.67 L/min. For consecutive 5-minute intervals the rates, rounded to nine decimals, are 0, 1.668107192, 1.343440056, 1.283871597, 0, 0 L/min; the interval duration times their sum is 21.477094 L. It adds 21.48 L, 11.0% below the constant rate. Its 0.1°C design margin sacrifices some water savings to protect the stated constraints. The remaining 6.06 L gap to the energy bound is not an optimality certificate.




Figure 3. Selected buffered 6-segment schedule against the best constant rate (top) and the range of cell temperatures outside the inlet jet zone (bottom); dotted lines mark the limits.

The unbuffered 12-segment candidate uses 19.35 L but reaches a spread of 1.531°C on the finest mesh, above the 1.5°C limit, so it is rejected. Among the tested buffered alternatives, the accepted 6-segment policy gives spreads of 1.400 , 1.395 and 1.416°C on the $8 \times 4 \times 3$, $12 \times 6 \times 4$ and $16 \times 8 \times 6$ meshes. Independently integrated continuous-time bounds also pass on each mesh. This establishes numerical feasibility for these tested networks, not mesh-independent physics or a global minimum.

7.2 What does staying close to the start temperature cost?

The task asks for a bath close to its initial temperature without wasting much water, which is a trade-off rather than a single optimum. Figure 4 repeats the search while the allowed fall below 40°C changes from 0.25 to 3°C ; the 1°C case is the baseline. With a fall of 2.5°C or more the stored heat suffices and no water is added. At 2°C the best constant rate adds 1.9 L, at 1.5°C 12.9 L, at 1°C 24.1 L and at 0.75°C 29.8 L. Between 2°C and 0.75°C each further 0.5°C of tolerance is therefore worth roughly 11 L, almost linearly. At 0.5°C the best constant rate adds 36.3 L; for a fall of 0.25°C no constant-rate policy was accepted, while the perfectly mixed bath would still need 30.7 L.

Figure 4 compares the spatial constant-rate search with the analytical well-mixed optimum and energy lower bound. The spatial curve joins accepted search results; it does not certify an optimum over all controls.

The curve is the practical answer to “how close is close enough”: below a 2.5°C fall every degree of tolerance bought back costs about the same, so the tolerance is worth choosing deliberately. The perfectly mixed curve lies below the spatial one at every tolerance, and the gap widens as the tolerance tightens, from 1.5 L at 2°C to 9.2 L at 0.75°C : the stricter the comfort requirement, the more uneven mixing costs. Only constant-rate policies are searched here; scheduled flow (Section




Figure 4. Water use versus temperature tolerance. Crosses denote unsuccessful searches, not zero water demand; the dotted line marks the baseline.

7.1) lowers the baseline point by about 11% at this one point. Schedules at other tolerances have not been optimized here.

7.3 Can a user follow the schedule?

We distinguish fixed tap bias from random segment errors. For the unbuffered 12-segment candidate, multiplying every rate by 0.9 or 1.1 crosses a limit by 0.076 and 0.126°C. Independently drawing each multiplier from a normal distribution of mean 1 and standard deviation 0.1, clipped below zero, gives median and 5th-percentile slacks of -0.070 and -0.217°C (200 draws, seed 11). Standard deviation 0.1 does not mean an error bounded within $\pm 10\%$. Negative slack is a violation, however small.

The error is not symmetric. The most sensitive segment is segment 4 (2.81 L/min), where raising the flow by 20% breaks the limits by 0.245°C and lowering it by 20% by 0.047°C: extra hot water overheats the inlet region and widens the spread. Over-delivery in a high-flow segment is the error to guard against. 1 of the 5 active segments tolerates a 20% error in either direction.

The price of a margin is measured on a 6-segment version. Requiring all three limits to hold with a buffer b gives 19.50 L at $b = 0$, 21.48 L (+10%) at 0.1°C, 23.81 L (+22%) at 0.2°C and 26.83 L (+38%) at 0.3°C. With the 0.1°C buffer, 74% of 200 draws with independent Gaussian multiplier standard deviation 0.1 stay within the limits. The 0.1°C policy is selected here for its independently checked numerical margin. Its 74% acceptance under the stated random error model does not make it a reliable manual prescription; a real bath needs temperature feedback and separate calibration.

7.4 Measurement ambiguity changes the decision

Mixing-specific candidates at known $D = 0.0007$ and 0.0013 m²/s use 27.00 and 20.30 L; the weaker-mixing candidate starts immediately. Both pass three-grid continuous checks. The finite local searches allow 2 L/min, against the baseline 3 L/min: this tests conditional actions, not equal-condition minimum water use. Nine fixed-policy cases also show that two probes can miss two sampled spatial violations. Measuring temperature is therefore not the same as identifying a usable model.

We test the decision consequence of that ambiguity with a separate synthetic calibration trial: a six-stage pulse (0, 0, 1.2, 0, 0, 0) L/min, five minutes per stage, observed every 30 s at fixed cold/hot cells 89 and 6. The geometry and uniform 40°C initial state are known. Each reading has an assumed $\pm 0.02^\circ\text{C}$ error and each probe an unknown constant $\pm 0.02^\circ\text{C}$ bias. These are hypothetical bounds, not instrument specifications.

For parameters θ and probe j , let $d_j(t; \theta) = T_j(t; \theta) - y_j(t)$. At observed times $\mathcal{T} = \{0, 30, \dots, 1800\}$ s, a model is compatible when a constant correction a_j exists such that

$$\min_{|a_j| \leq 0.02} \max_{t \in \mathcal{T}} |d_j(t; \theta) - a_j| \leq 0.02 \quad (j = 1, 2). \quad (18)$$

The minimax correction is the midpoint of the maximum and minimum discrepancy, clipped to the bias bound. In the observation equation $y_j = T_j + b_j + e_j$, the physical bias is $b_j = -a_j$. A varying trace from -0.04 to $+0.04^\circ\text{C}$ fails this test even though each reading lies within $\pm 0.04^\circ\text{C}$: the offset must remain constant.

We enumerate 7 mixing values from 0.00085 to 0.00115 m²/s, nine surface coefficients from 20 to 30 and nine body coefficients from 15 to 35 W/(m² K), and five delivery multipliers from 0.95 to 1.05. Of 2835 models, 178 fit these synthetic observations. The original six-stage policy fails the five-second sampled limits in 4 compatible models. The finite set is a decision ambiguity set, not a confidence region or probability distribution.

Table 8 shows a replacement candidate: optimize ten extreme or failing models with the baseline 0.1°C design target and additional 0.03°C floor reserve, then check all 178. The local search has a 2 L/min bound and does not establish global optimality. Its three-grid conditional floating-point envelopes pass, but the tightest spread bound is 1.496828°C: matching the design target does not give every model a 0.1°C final margin.

The replacement uses 9.33% more commanded water than the nominal 21.48 L policy. This is a finite-set feasibility repair, not a saving under a fair uncertainty comparison. The pulse additionally uses 6 L (5.7–6.3 L delivered), and both experiments start at 40°C; fill/reset costs are not modeled. Reusing calibration across n unchanged baths amortizes only the pulse to $6/n$ L. No online feedback, continuous parameter coverage or real-bath reliability follows from this experiment.

Table 8. Control under measurement ambiguity

Quantity	Conditional result
Six successive rates (L/min)	0.892820, 1.473786, 1.179268, 1.150408, 0.000000, 0.000000
Commanded / delivered water	23.48 L / 22.31–24.66 L
Minimum / maximum envelope	39.1297 / 40.9238°C
Maximum spread envelope	1.496828°C; 534 model-grid checks
Separate pulse + control	29.48 L commanded; reset cost unknown

8 Separate geometry, size and body effects

The scenarios in Table 9 distinguish changes in water storage from changes in body displacement and heat exchange; each row uses its own reconstructed geometry.

Table 9. Geometry, size and body scenarios

Change from baseline	Water (L)	Interpretation
Shallow Wide	No candidate	Same envelope volume
Deep Narrow	17.48	Same envelope volume
Small Bath	27.05	Lower water depth
Large Bath	20.39	Higher water depth
Deep Bath	13.05	Depth 0.35 m, like a bath filled to the shoulders
Larger Body	28.11	More displacement/contact
Long Body	23.85	Shape only; fixed volume/area
Warmer Skin	21.08	Skin at 36°C
Cooler Skin	27.18	Skin at 32°C

At equal volume, the wide, shallow tub has a larger exposed top area and a changed diffusion length; it has no accepted candidate. The deep, narrow tub needs 17.48 L against 24.14 L, because a deep, narrow envelope reduces top-area heat loss, although a real deep bath may stratify, which Section 9 tests with a separate vertical-mixing scenario. Geometry is therefore more than an interchangeable cooling coefficient.

Changing depth at fixed length and width changes thermal storage and shell area. Greater water volume slows cooling, which can lower added water over a short fixed horizon: the large bath needs 20.39 L, while the small bath needs 27.05 L. The 0.23 m baseline depth is shallow, and its surface-to-volume ratio is about 1.7 times that of a bath filled to the shoulders; a 0.35 m depth needs 13.05 L, so the baseline overstates the surface-loss share of a deeper bath. That does not mean a larger bath conserves total water: filling it initially uses more. J counts replenishment only, and this distinction prevents a misleading conservation claim.

The larger-body scenario simultaneously displaces more water and increases assumed contact area; it is a combined size scenario and needs 28.11 L. The long-body scenario isolates the spatial distribution by keeping total body volume and contact area fixed. Warmer skin reduces the temper-

ature difference driving body heat loss. These changes quantify conditional dependence; they do not establish measured human heat transfer.

To transfer the model to a real tub, measure water volume after entry, submerged contact geometry, temperatures at several depths and distances, and no-inlet cooling. Shape and temperature should not be inferred from a single mean cooling curve.

9 Motion and a bubble layer can change the policy

The outcomes in Table 10 show why mixing and boundary heat loss must be varied separately: stronger transport can help, while added surface loss can offset the gain.

Table 10. Mixing, surface, comfort and supply scenarios

Scenario	Added water (L)	Search outcome
Weak Mixing	—	None accepted
Strong Mixing	19.40	Accepted
Moving With Added Surface Loss	23.04	Accepted
Stratified	24.42	Accepted
Convective Surface Layer	24.10	Accepted
Foam	10.37	Accepted
Tight Comfort	36.32	Accepted
Loose Comfort	1.87	Accepted
Low Loss	8.60	Accepted
High Loss	—	None accepted
Cool Supply	44.80	Accepted

Increasing D reduces the gradient created by the localized inlet. The strong-mixing scenario uses 19.40 L against 24.14 L, but motion may also increase heat loss. When D is increased together with a 20% increase in the surface coefficient (30 instead of 25 W/(m² K)), the need rises to 23.04 L and most of the benefit disappears. The comparison deliberately separates transport improvement from its possible boundary cost.

Stratification is tested separately. A hot inlet layer is lighter than the water below, and the stable density gradient suppresses vertical mixing. With a cell height of 0.077 m, a vertical temperature difference of 1 K, a velocity scale of 0.02 m/s and an expansion coefficient of about 3.8×10^{-4} per K, the gradient Richardson number is $Ri = g\beta \Delta T \ell / u^2 \approx 0.7$. The Munk–Anderson stability function for scalars, $(1 + 3.33 Ri)^{-3/2}$ [8], then reduces vertical diffusivity to about 0.15 of its neutral value. We test a vertical-to-horizontal ratio of 0.2. The best constant rate becomes 24.42 L against 24.14 L, but the maximum spread rises from 0.97 to 1.07°C against the 1.50°C limit. Stratification therefore costs little water here and uses 19% of the remaining uniformity margin; the ratio is an order-of-magnitude scenario and the Richardson estimate uses assumed velocity and temperature scales. The stable layer is an upper-bound picture: evaporative cooling at the surface can instead

drive convection and raise vertical mixing, so a vertical-to-horizontal ratio of 3 is also tested, and it needs 24.10 L.

The foam scenario multiplies only the effective surface coefficient by 0.4. With all other inputs fixed, the best constant rate falls to 10.37 L. The assumed 60% reduction is a scenario, not an experimentally established property of bubble-bath additive. If the layer breaks up or motion raises evaporation, this result must be recomputed.

No accepted candidate exists for: Weak Mixing, Shallow Wide, High Loss. They are reported as such, not hidden in a favorable average. In the weak-mixing, high-loss and wide-shallow cases, flows large enough to keep every cell above the floor already push the spread above its 1.5°C limit (checked at constant flows from 0.8 to 2.5 L/min), but these sampled candidates cannot establish that every constant flow fails. No constant-rate policy was accepted by the documented search in those scenarios. Their outcomes identify where the restricted strategy must change. Potential responses include improved circulation, a different inlet path, a shorter bath, or a different comfort tolerance. The paper does not certify which response is safe or optimal without corresponding physical evidence.

10 Validation: independent evidence and remaining error

All 17 recorded checks pass. Independent heat-flow arithmetic sums environmental losses, body exchange and local-temperature overflow; the largest instantaneous residual is 2.46×10^{-10} W. An adaptive RK45 integration with independently assembled right-hand side agrees with the archived matrix-exponential trajectory within 3.12×10^{-07} °C. This tests numerical implementation, not a second physical model.

A constructed uniform-loss problem reproduces its analytic exponential cooling solution. Internal exchange cancels globally, and the stream has exactly one overflow sink. Positive off-diagonal dynamics and nonpositive row sums support physical bounds and a contractive derivative estimate for each constant-flow segment.

Half-second replay is supplemented by a derivative bound: if $|dT_i/dt| \leq L$, a point lies at most 0.25 seconds from a sample, so temperatures differ by at most $L/4$ and spread by at most $L/2$. With an integration allowance and no relaxation of the physical limits, the resulting lower envelope is 39.0144°C, upper 40.3340°C, and spread upper bound 0.9980°C. This is a conditional floating-point envelope, not interval-arithmetic certification.

The accepted schedule uses an independently assembled right-hand side and adaptive RK45, restarted at every flow switch. Nonnegative off-diagonal transport and nonpositive row sums give an infinity-norm contraction bound for the derivative within each segment. One-second samples alone leave some intervals unresolved; adding RK45 step endpoints and bisecting unresolved intervals once on the finest mesh gives bounds of 39.1297°C for the minimum, 40.8304°C for the maximum and 1.4964°C for the spread. These pass the physical limits without threshold relaxation. A 2×10^{-6} °C integration allowance is included; this floating-point check is conditional on the network, not interval arithmetic or experimental validation. The 17 baseline checks are separate.

11 Resolution, uncertainty and transfer to a real bath

The three-grid replay in Table 11 checks the quantities constrained outside the fixed jet zone and separately reports the inlet-cell peak, which is not covered by that ceiling.

Table 11. Mesh replay: constrained temperatures

Mesh (cells)	Max outside jet zone	Spread outside zone	Inlet-cell max
$8 \times 4 \times 3$ (96)	40.318°C	0.967°C	40.553°C
$12 \times 6 \times 4$ (288)	40.315°C	0.960°C	40.757°C
$16 \times 8 \times 6$ (768)	40.330°C	0.973°C	40.995°C

The same constant-rate policy is replayed on three meshes. The ceiling and spread quantities outside the jet zone agree within 0.015°C and 0.013°C, and the tested constant-rate trajectory stays within the comfort limits on these meshes. The inlet-cell maximum rises from 40.55 to 41.00°C with refinement; that is why the cell is excluded rather than constrained. The selected water amounts on the $8 \times 4 \times 3$ and $12 \times 6 \times 4$ meshes are 24.142 and 24.117 L (0.11% apart). Three meshes are a diagnostic, not an asymptotic convergence-order study.

Parameter uncertainty concerns heat loss, body exchange and mixing. Structural uncertainty concerns the prescribed surface stream, fixed skin reservoir, homogenized body and constant effective D . These are different errors. The scenario table explores parameter dependence and one coupled motion effect; it is not a probability distribution or confidence interval.

Range analysis. Section 2 gives ranges for six inputs. 64 scrambled Sobol points cover them (D log-uniform), and the constant-rate search runs at each. 41 points have an accepted candidate. Among those, the required water is 12.5, 17.0, 22.0, 26.9 and 30.9 L at the 5th, 25th, 50th, 75th and 95th percentiles. The remaining 23 points have none, mostly because of weak mixing (below). The baseline of 24.1 L sits in the upper half of this distribution.

The draws in Figure 5 show how surface heat loss changes water demand within the tested ranges, while the rejected draws expose conditions where the search finds no acceptable policy.

Spearman rank correlations with water, over accepted draws: surface coefficient +0.63, room temperature -0.47, body coefficient +0.47, skin temperature -0.27, shell coefficient +0.15, mixing diffusivity -0.15.

The correlations use accepted draws only, so they say nothing about which draws have an accepted policy at all. In this finite search, acceptance differs most strongly with the mixing diffusivity: its median is 1.4×10^{-03} m²/s for draws with an accepted policy and 4.5×10^{-04} for draws without (Mann–Whitney $p < 1 \times 10^{-4}$), while no association is detected for the other five inputs in this sample (smallest $p = 0.47$). Weak mixing leaves the far water too cold unless the flow is raised so far that the inlet region breaks the spread limit. The draws are uniform over ranges that we chose; they are not a probability distribution for any real tub. The main message: among accepted candidates, the loss coefficients and room temperature affect the water needed, while the mixing closure strongly affects search acceptance, which is why Section 2 ties the coefficients to outside




Figure 5. Required water against the surface coefficient for each Sobol draw. Crosses at the top mark draws with no accepted policy; the vertical line is the baseline.

evidence and Section 9 reports weak mixing separately.

A calibration plan should begin with an inlet-free cooling experiment, recording water volume and room conditions. Multiple temperature probes distinguish aggregate heat loss from mixing. A separate inlet experiment records flow and inlet/outlet temperatures, allowing overflow energy to be checked. Motion and foam need matched trials because both can alter transfer coefficients. Reserve an entire experiment for prediction checks after fitting. A single mean-temperature series cannot separate the surface, shell and body coefficients, so only calibrated and independently tested parameters would support advice for a particular bath.

11.1 Do the policies survive a different body or flow path?

We replay both archived policies on the 96-cell network under two structural alternatives and their combination, without re-optimizing. A finite contact reservoir replaces fixed skin: $C_b \dot{B} = \sum_i H_{b,i}(T_i - B)$, with 253.31 kJ/K and initial $B = 34^\circ\text{C}$. This tests storage, not physiology; thermoregulation is omitted. The alternative face-connected stream descends at the inlet, crosses the bottom, then rises to the unchanged overflow. Geometry, mixing, losses and the jet zone remain fixed.

Table 12 compares the unchanged scheduled policy under the baseline and limited structural alternatives.

Both policies meet the sampled limits in all four structures. Body storage changes the scheduled minimum by $+0.039^\circ\text{C}$; the deep route changes its spread by $+0.023^\circ\text{C}$. The effective body node ends near 35.05°C . Thus these limited changes preserve sampled feasibility, not optimality or the 11% saving after re-optimization.

Independent heat-flow arithmetic and RK45 restart at each switch. Sampling at 1 and 0.5 s changes extrema by at most $4.83 \times 10^{-05}^\circ\text{C}$; maximum energy residuals are 1.63×10^{-10} W in-

Table 12. Fixed-policy structural replay

Scheduled-policy structure	Min (°C)	Max (°C)	Spread (°C)
Baseline	39.130	40.774	1.400
Finite contact reservoir	39.169	40.788	1.397
Deep transport path	39.130	40.794	1.423
Both alternatives	39.169	40.808	1.421

stantaneously and 4.03×10^{-08} J cumulatively. These alternatives lack the baseline's three-grid continuous-time check. The fixed-skin 15.41 L bound does not transfer unchanged to a warming body reservoir. No alternative is physically calibrated.

12 Conclusions and a policy with clear scope

The model supports five linked findings. First, tolerating a modest temperature decline lets stored heat replace some inlet water. Second, spatial transport changes the ideal mixed recommendation: the best constant rate starts at once, and a schedule that withholds water at both ends needs 11% less. Third, among accepted candidates in our sensitivity sample, the surface and body loss coefficients are important determinants of water use: from 12 to 31 L across the scenario ranges. Fourth, improving mixing helps only to the extent that it does not create offsetting boundary loss or hot-water short circuit. Fifth, the accepted timing changes across our tested mixing conditions: the $D = 0.0007$ m²/s candidate starts immediately, while the baseline waits. Two fixed temperature readings can miss marginal spatial violations and do not by themselves validate a feedback policy.

For the stated baseline, the best constant trickle is 0.805 L/min for 24.14 L over 30 minutes; the 6-segment schedule uses 21.48 L. The scheduled result includes a 0.1°C design margin and three-grid continuous-time checks; the floor covers every cell, while the ceiling and spread exclude the inlet zone. Neither is a universal faucet prescription or a proved global optimum over all time-varying controls.

Strengths. Exact network conservation, independent integration, proved benchmark and lower bound connect mechanism with decision. Scheduling reduces the nominal bound gap from 8.7 to 6.1 L; physical accuracy remains untested.

Limitations. Local policy search, assumed mixing and jet zone, omitted shell storage, fixed skin and unmodeled flow physics limit transfer. Section 11.1 tests only selected structural alternatives. Preference limits and cell averages do not establish inlet-jet safety.

Parameter estimation need not be unique to inform a decision: propagate the remaining compatible models into policy checks. Here the nominal six-stage policy fails four of 178 synthetic-compatible models; the 23.48 L replacement passes three-grid envelopes but costs more, and its separate pulse adds 6 L before unknown reset costs. The finite set is not a reality guarantee. Practical rates still require measured cooling and mixing; geometry or motion changes require reassessment. The following guide translates these principles.

13 A warmer bath, with less replacement water

A guide for the person in the bathtub

Decide what “warm enough” means. A bath need not stay at exactly its starting temperature to remain acceptable. In our example, allowing a 1°C fall still needs about 24 L of replacement water, while allowing a 2°C fall needs only 1.9 L. Your actual comfort and health needs must determine the acceptable range.

Distribute the warmth before turning up the tap. The water near the faucet can warm while the far end remains cool. Gentle movement can help redistribute heat. More vigorous motion is not automatically better: it may also increase heat loss from the surface.

Time the inflow for this scenario. The verified 6-stage schedule waits 5 minutes, adds hot water during the next 15 minutes, then stops for 10 minutes. It uses 21.48 L against 24.14 L for a constant trickle. This timing was checked for the stated tub, heat losses and mixing conditions only. A limited body-storage and flow-path replay preserves the example’s sampled limits, but does not establish transfer to other baths. Other tubs or body positions need a new calculation; if the water is not mixing or the thermometer disagrees, the schedule is not transferable.

Leave room for error. The lowest-water candidate failed our refined calculation. Even with a margin, only 74% of trials with independent Gaussian tap multipliers of standard deviation 0.1 stayed within the limits. This is a model experiment, not your success probability. More water is not automatically better: a strong burst can overheat the inlet region.

Check more than one location, and do not assume two readings cover the bath. Our fixed-probe study misses spatial violations even without sensor noise. Reconsider timing when cold or hot locations move. The separate calibration trial adds 6 L before filling or resetting costs; it is not free feedback during the bath. Cell-average temperature does not establish hot-jet safety.

Remember where the added water goes. Once the tub is full, extra water leaves through the overflow. Hot water that runs along the surface to the overflow may escape before its warmth reaches you. Increasing the stream can waste heat without fixing unevenness.

A bubble layer may help, but the amount is uncertain. An intact insulating surface layer can reduce heat loss. Our calculation explores such a layer; it does not measure the effect of a particular additive. If the layer disappears or the bath behaves differently, reassess rather than relying on the example’s saving.

The practical rule: use the bath’s stored warmth, keep temperature distributed, and add only as much water as your actual conditions require. The numerical example explains these tradeoffs; it is not a tested bathing or safety standard.

14 References and reproducible algorithm

- [1] COMAP. 2016 MCM Problem A: A Hot Bath. 2016. Official problem PDF: contest.comap.com/undergraduate/contests/mcm/contests/2016/problems/2016_MCM_Problem_A.pdf.
- [2] U.S. Department of Energy. DOE Fundamentals Handbook: Thermodynamics, Heat Transfer, and Fluid Flow. DOE-HDBK-1012/2-92, June 1992. Convection heat transfer, Eq. (2-9).
- [3] IAPWS. Revised Release on the IAPWS Formulation 1995 for the Thermodynamic Properties of Ordinary Water Substance for General and Scientific Use. R6-95(2018). iapws.org/technical-guidance/release/IAPWS-95.
- [4] SciPy developers. `scipy.sparse.linalg.expm_multiply` and `scipy.integrate.solve_ivp`, API documentation. docs.scipy.org/doc/scipy/reference/. Accessed 7 October 2026.
- [5] Bergman, T. L., Lavine, A. S., Incropera, F. P., DeWitt, D. P. Fundamentals of Heat and Mass Transfer, 7th ed. Wiley, 2011. Natural convection above a heated horizontal surface (characteristic length A/P); heat and mass transfer analogy.
- [6] Menzies, C., Clarke, N., Steward, C. J., Thake, C. D., Pugh, C. J. A., Cullen, T. Vascular, inflammatory and perceptual responses to hot water immersion: impacts of water depth and temperature in young healthy adults. *Experimental Physiology*, 2025. doi:10.1113/EP092761.
- [7] OpenAI, Codex (GPT-6-based assistant), and Anthropic, Claude Sonnet 5.5 in Claude Code. Used on 7 October 2026 for modeling, code, validation and report composition; exact builds not independently established. See Report on Use of AI.
- [8] Munk, W. H., Anderson, E. R. Notes on a theory of the thermocline. *Journal of Marine Research*, 7(3), 276–295, 1948.

Algorithm and computational record

1. Form cell volumes and capacities; reject excessive body occupancy. Assemble symmetric face exchange, environmental/body conductances and the conservative stream. 2. Compute the mixed closed forms and energy lower bound. 3. Evaluate delayed constant-flow candidates with matrix exponentials and refine the first bracketed lower-temperature crossing. Reject ceiling or spread violations outside the jet zone. 4. Rank accepted candidates by added litres. 5. Independently integrate and check balances, analytic limits, parameter arithmetic and constraints; replay a finer mesh and perturb physical scenarios. 6. Optimize piecewise-constant and buffered schedules, then accept only candidates whose independent three-grid continuous bounds pass. 7. Rerun the constant-rate search at Sobol draws over the literature ranges. Independent RK45 checks use relative tolerance 2×10^{-9} , absolute $2 \times 10^{-10}^\circ\text{C}$ and steps of at most 5 s. AI assistance is acknowledged in [7].

Appendix A. Scenario definitions

Each row changes only the listed baseline inputs. All unspecified inputs remain as stated in Section 2; geometry and conductance are then rebuilt. The equal-volume aspect-ratio cases retain the envelope volume of 0.22425 m^3 . Widths s describe a smooth displacement distribution, not measured

anatomical semiaxes.

Accepted candidates are independently replayed on their scenario network at intervals of no more than five seconds. These scenario checks are sampled checks; continuous-time envelopes are reported for the baseline constant rate and the selected scheduled policy; the latter uses additional interval refinement when necessary. “None accepted” refers to the documented finite search and is not an infeasibility theorem.

The inputs in Table 13 and Table 14 make the scenario comparisons reproducible: only the listed quantities change from the baseline.

Table 13. Scenario definitions

Scenario	Changed input(s)
Weak Mixing	$D \text{ (m}^2/\text{s)} = 0.0003$
Strong Mixing	$D \text{ (m}^2/\text{s)} = 0.003$
Moving With Added Surface Loss	$D \text{ (m}^2/\text{s)} = 0.003$; surface coefficient ($\text{W}/\text{m}^2\text{K}$) = 30
Loose Comfort	lower limit ($^{\circ}\text{C}$) = 38
Tight Comfort	lower limit ($^{\circ}\text{C}$) = 39.5
Foam	surface multiplier = 0.4
Shallow Wide	length (m) = 1.7; width (m) = 0.75; depth (m) = 0.175882
Deep Narrow	length (m) = 1.3; width (m) = 0.6; depth (m) = 0.2875
Small Bath	depth (m) = 0.2
Large Bath	depth (m) = 0.27

Table 14. Scenario definitions (continued)

Scenario	Changed input(s)
Deep Bath	depth (m) = 0.35
Larger Body	body displacement (m^3) = 0.075; body contact area (m^2) = 1.35
Long Body	body widths s (m) = [0.48, 0.12, 0.1]
Warmer Skin	skin temperature ($^{\circ}\text{C}$) = 36
Cooler Skin	skin temperature ($^{\circ}\text{C}$) = 32
High Loss	surface coefficient ($\text{W}/\text{m}^2\text{K}$) = 36; wall coefficient ($\text{W}/\text{m}^2\text{K}$) = 8.5; body coefficient ($\text{W}/\text{m}^2\text{K}$) = 40
Low Loss	surface coefficient ($\text{W}/\text{m}^2\text{K}$) = 17; wall coefficient ($\text{W}/\text{m}^2\text{K}$) = 4.5; body coefficient ($\text{W}/\text{m}^2\text{K}$) = 12
Stratified	vertical/horizontal D ratio = 0.2
Convective Surface Layer	vertical/horizontal D ratio = 3
Cool Supply	inlet temperature ($^{\circ}\text{C}$) = 45

Report on Use of AI

Tool and scope

OpenAI Codex, a GPT-6-based assistant, and Anthropic Claude Sonnet 5.5 through Claude Code assisted development on 7–9 October 2026; exact builds were not independently established. They formulated models, wrote code, checked calculations, drew figures and composed the report. The earlier final version was sealed before a later comparative-learning revision. That revision read selected portions of same-problem Outstanding papers 44845 and 54164 and COMAP's judges' commentary; it must not be described as an unseen-problem trial. Their data and figures were not copied. An assistant sub-agent wrote the independent-RHS validation routine and another checked fixed probes; this is AI-assisted review, not independent human verification.

Task

The task was to select a classic historical MCM problem and build a rigorous, professional case from it. The conversation wording is not reproduced in this report or in the development record.

Outputs, corrections and verification

On 8 October, independent integration rejected the 19.35 L schedule and accepted a buffered alternative; later fixed-policy replays checked limited body storage and flow-path changes. On 9 October, the baseline values were bound to archived results without repeating that optimization. Earlier coefficient and inlet-zone corrections are recorded in the accompanying development history. References [1]–[4] were checked at the linked sources and [6] in its open manuscript; [5] and [8] were not opened, so their relations were checked indirectly. The AI produced the code, derivations, figures and report; all reported numerical results came from actual calculations.

The later control study reused the accepted baseline, separately optimized schedules at two specified diffusivities, independently checked three-grid continuous envelopes, and replayed nine fixed-policy cases with fixed temperature probes. A separate synthetic pulse screened 2,835 parameter combinations, retained 178 measurement-compatible models and checked a replacement schedule on three grids. Calibration and reset costs are distinguished from control water. The original baseline optimization was not rerun. The accompanying study source and hash-bound receipt retain finite search scope and sensor omissions. Additional exploratory feedback candidates were not adopted as recommendations.

Record limitations

A full exported interaction transcript and every intermediate AI output were not available in this artifact. The task description, tool identification, scope, code artifacts and numerical receipts are retained in the accompanying development record; this report does not invent a transcript. No independent human review or physical bath experiment is claimed.