

Problem Chosen
A

MCM/ICM
Summary Sheet

Team Control Number
7391856

A Hot Bath: Conserving Water Without Losing Uniformity

Summary

Maintaining a warm bath is a coupled problem of heat loss, replenishment and transport. A hot inlet can improve the mean temperature while leaving distant water cool and sending useful heat directly to the overflow. We therefore minimize added water subject to explicit limits on every modeled cell, rather than optimizing an average alone.

We prove a coast-then-hold policy optimal for a well-mixed bath, then use a three-dimensional finite-volume thermal network with coefficients anchored in standard correlations and a published immersion study (Section 2). The best constant rate is compared with an optimized piecewise-constant schedule; neither is proved optimal among all controls.

Recommendation. In the baseline tub, use the stored heat at both ends of the bath and run the tap in the middle: no hot water for the first 2.5 minutes or the last 10 minutes, and 0.8–1.6 L/min in between, which adds 19.8 L over 30 minutes. A constant trickle needs 24.1 L (0.80 L/min), 22% more. The model leads to this schedule because every cell must stay within limits, so heat has to reach the coldest cell, not just raise the mean.

In a 164.25 L water-volume scenario lasting 30 minutes, with an initial temperature of 40°C and a 39°C lower limit, the best constant-rate policy adds **24.14 L** at 0.805 L/min from the start, and a 12-segment schedule found by constrained optimization needs **19.77 L**, 18% less. The well-mixed optimum is 16.01 L, while an independent energy argument gives a 15.41 L lower bound for the spatial problem. A finer mesh changes the constant-rate result by 0.11%.

The answer depends strongly on the surface and body loss coefficients. Across 64 Sobol draws over literature-based ranges, 37 had an accepted constant-rate policy, with required water from 12 to 31 L (5th–95th percentile, median 21 L). Weak mixing yields no accepted candidate under the stated temperature limits; stronger mixing saves water, but added surface loss can reverse that benefit. In the assumed foam scenario replenishment falls to 10.4 L. These are conditional calculations, not measurements or product claims.

Independent integration, analytic limits, energy accounting and a between-samples envelope support the results: 17 model checks and 2 for the schedule. A one-page user guide and the limitations follow; AI assistance is disclosed in the references and appended report.

1. Define the decision before optimizing

The task is to preserve both warmth and spatial uniformity in an overflowing, unheated tub, and to examine geometry, the bather and motion, and a bubble-bath layer [1]. The report separates physical requirements from preference assumptions. There is no supplied temperature record or measured heat-transfer coefficient to fit.

The decision variables are an inlet flow rate and the time at which a constant trickle begins. The tub is already full: added water displaces an equal volume through the overflow. The horizon is 1,800 s. Our baseline accepts cell averages between 39°C and 41°C, with an instantaneous spread of at most 1.5°C. These choices operationalize comfort; they are not medical limits or numbers specified by the problem.

$$J = 1000 \int_0^{t_f} q(t) dt \quad (1)$$

$$T_{\min} \leq T_i(t) \leq T_{\max}, \quad \max_i T_i(t) - \min_i T_i(t) \leq \Delta \quad (2)$$

Here q is in m^3/s and J is in litres. We prioritize the least water within the constraints, rather than assigning arbitrary weights to unlike units. Tightening the temperature tolerance is a separate scenario. Zero flow is admitted: if the initial stored heat suffices, using no added water is globally water-minimal.

Requirement	Where answered
Temperature in space and time	Sections 3, 5–7; Figures 1–2
Water-use strategy and its scope	Sections 4, 6–7 and 12
Tub/body geometry, size and temperature	Section 8
Motion and bubble-bath additive	Section 9
Validation and sensitivity	Sections 8–11
One-page non-technical explanation	Section 13

Spatial temperature means a control-volume average. It does not bound the unresolved temperature of a faucet jet or a skin-contact film. This distinction determines which practical conclusions the simulation can support.

2. Physical assumptions and scenario inputs

A constant-density, constant-heat-capacity liquid stores sensible heat. The exchange law is proportional to temperature difference; coefficients depend on the situation and normally require empirical calibration [2]. Rounded water-property constants are used as approximations, not as an evaluated IAPWS property-table dataset [3].

Quantity	Baseline / units	Status
Tub $L \times W \times H$	$1.50 \times 0.65 \times 0.23$ m	Assumed geometry
Displaced body volume / area	$0.060 \text{ m}^3 / 1.15 \text{ m}^2$	Assumed
Water density / heat capacity	$1000 \text{ kg/m}^3 / 4180 \text{ J/(kg K)}$	Rounded constants
Air / skin / inlet temperature	$22 / 34 / 50^\circ\text{C}$	Fixed reservoirs
Surface / wall / body coefficient	$25 / 6.5 / 25 \text{ W/(m}^2 \text{ K)}$	Literature-anchored (next page)
Mixing diffusivity D	$0.001 \text{ m}^2/\text{s}$	Uncalibrated closure; range tested
Initial / lower / upper limit	$40 / 39 / 41^\circ\text{C}$	Preference scenario
Time / allowed spread	$1800 \text{ s} / 1.5^\circ\text{C}$	Preference scenario

The surface coefficient aggregates convection, linearized radiation and evaporation over a narrow temperature range. The shell coefficient aggregates water-side transfer, wall conduction and external loss to room air. This avoids adding an evaporation term twice; it does not predict humidity or foam chemistry separately. Evaporation-related volume change is neglected relative to the water volume; the boundary coefficient models its heat effect only.

Skin temperature is held fixed as an effective reservoir. A 30-minute bath can change skin temperature; the fixed value is therefore varied, not claimed to resolve thermoregulation. Bath motion is represented through D , with an additional surface-loss scenario. Coefficients cannot be identified from this problem statement alone.

2.1 Where the coefficients come from

No temperature record exists to fit, so each coefficient is anchored in a standard relation or a published measurement and given a range. Surface loss sums natural convection above a hot horizontal surface ($Nu = 0.15 Ra^{(1/3)}$ for $10^7 < Ra < 10^{11}$, length A/P [6]; here $Ra = 1.9 \times 10^7$), linearized radiation (emissivity 0.96) and evaporation by the Lewis analogy. For open water at 40°C in 22°C air at 50% humidity the parts are 4.7, 6.1 and 24.8 W/(m² K) (0.66 kg/m² h evaporated), 35.6 in total. A bather covers part of the surface; with an assumed exposed fraction of 0.7 the central value is 24.9, and we use 25 within the range 17–37.

The shell coefficient is a series resistance: a 5 mm shell (0.19 W/(m K)) with film coefficients of 300 inside and 8 W/(m² K) outside gives 6.5, and 4.7–8.4 for other thicknesses. These are typical engineering values, not measurements of a particular tub. The body coefficient is anchored by a measurement: Menzies et al. report a rectal-temperature rise of $0.9 \pm 0.3^\circ\text{C}$ after 30 minutes in 40°C water to the shoulders [7]. Assuming a 75 kg body, $c = 3470 \text{ J/(kg K)}$ and a mean body rise of 1–2°C, the average heat uptake is 145–289 W, equal to 21–42 W/(m² K) with skin held at 34°C. We use 25, because a fixed skin node overstates the driving difference later in the bath. The mass and the mean rise are assumptions, so this is an order-of-magnitude anchor, not a calibration.

Mixing has no anchor: $D = 0.001 \text{ m}^2/\text{s}$ is a scenario inside the range 0.0003–0.003. The 39–41°C window is a preference; immersion studies use 40–42°C water [7], which is not a safety standard. A 40.5°C upper limit was tried first: the coarse mesh accepted a 24.51 L constant-rate policy, but the finer mesh accepted none, because the inlet-cell average rises under refinement. The limit was set to 41°C, where both meshes agree.

Coefficient	Relation or anchor	Derived value	Used (range tested)
Surface, W/(m ² K)	Convection + radiation + evaporation [6]	35.6 open; 24.9 at 70%	25 (17–37)
Shell, W/(m ² K)	Series resistance, typical shell values	6.5 (4.7–8.4)	6.5 (4.5–8.5)
Body, W/(m ² K)	Uptake implied by core rise [7]	21–42 at fixed skin	25 (12–40)
Mixing D, m ² /s	No anchor; scenario	—	0.001 (0.0003–0.003)

3. A transparent well-mixed benchmark

Let C be total water heat capacity, H_a the air/shell conductance and H_b the body conductance. A well-mixed overflow has the same temperature as the bath. Integrating the physical energy balance gives

$$C\dot{T} = H_a(T_a - T) + H_b(T_b - T) + \rho c_p q(T_{\text{in}} - T) \quad (3)$$

$$C = \rho c_p(LWH - V_b), \quad H_b = h_b A_b \quad (4)$$

$$H_a = h_s f LW + h_w \{LW + 2H(L + W)\} \quad (5)$$

The foam multiplier f equals one without a layer. Setting $q = 0$, define $H = H_a + H_b$ and $T_e = (H_a T_a + H_b T_b)/H$. The cooling solution and first time to reach the lower limit are

$$T(t) = T_e + (T_0 - T_e)e^{-Ht/C} \quad (6)$$

$$t_c = \frac{C}{H} \log \frac{T_0 - T_e}{T_{\text{min}} - T_e} \quad (7)$$

Baseline conductances are $H_a = 37.14$ W/K and $H_b = 28.75$ W/K, and $C = 686,565$ J/K. Thus $T_e = 27.24^\circ\text{C}$ and $t_c = 14.17$ min.

A heat-loss model must approach an environmental equilibrium rather than zero Celsius. This analytic boundary also supplies a direct check on signs and units. The benchmark describes mixing perfectly; the spatial model will test the cost of departing from that assumption.

4. What can be proved about water use?

Let $\ell(T) = H_a(T - T_a) + H_b(T - T_b)$ be the heat loss of the mixed bath, and let the inlet be hotter than every temperature considered. **Proposition 1.** Coasting to the lower limit and then holding it uses the least water, provided the holding rate is within the faucet bound. The proof has four steps.

Step 1, an identity. Dividing the balance by $T_{\text{in}} - T$ and integrating over the horizon gives

$$\int q \, dt = \frac{C}{\rho c_p} \log \frac{T_{\text{in}} - T_0}{T_{\text{in}} - T_f} + \int \frac{\ell(T)}{\rho c_p(T_{\text{in}} - T)} \, dt \quad (8)$$

Step 2, monotonicity. The integrand $\ell(T)/(T_{\text{in}} - T)$ has derivative $[H_a(T_{\text{in}} - T_a) + H_b(T_{\text{in}} - T_b)]/(T_{\text{in}} - T)^2 > 0$, so it increases with T , and the logarithmic term increases with the final temperature. Water use is therefore an increasing functional of the temperature path.

Step 3, the lowest admissible path. Because $q \geq 0$, comparison with the no-flow solution gives $T(t) \geq T_c(t)$, the cooling curve of Section 3, and feasibility requires $T(t) \geq T_{\min}$. The pointwise lowest admissible path is $\max(T_c(t), T_{\min})$: coast until T_c reaches $T_{\min}$, then hold.

Step 4, attainability. Holding $T = T_{\min}$ means $\dot{T} = 0$, which needs the rate

$$q_{\text{hold}} = \frac{\ell(T_{\min})}{\rho c_p (T_{\text{in}} - T_{\min})} \quad (9)$$

When this rate is within the bound, the path is feasible, and by Step 2 it is optimal. $\square$ For the baseline it adds 16.01 L: wait 14.17 min, then supply 1.012 L/min. This is a proved optimum of the mixed model, not a guarantee for a spatially nonuniform tub.

Proposition 2 (energy lower bound). In the spatial model, suppose every cell satisfies $T_i(t) \geq T_{\min}$. Summing the cell balances gives $\sum C_i T_i(t_f) - C T_0 = -[L \, dt + \rho c_p] q (T_{\text{in}} - T_{\text{out}}) dt$, with L the total loss. Each loss term increases with its cell temperature, and H_a and H_b are the sums of the cell conductances, so $L(t) \geq \ell(T_{\min})$. The outlet cell satisfies $T_{\text{out}} \geq T_{\min}$, so $T_{\text{in}} - T_{\text{out}} \leq T_{\text{in}} - T_{\min}$, and the left side is at least $C(T_{\min} - T_0)$. Combining these three facts,

$$J \geq \max \left(0, \frac{1000[\ell(T_{\min})t_f - C(T_0 - T_{\min})]}{\rho c_p (T_{\text{in}} - T_{\min})} \right) \quad (10)$$

The resulting 15.41 L is a genuine conditional lower bound, valid for any control, but it is not tight: the initial water is hotter than the floor and loses more heat, and limited mixing adds further cost. Section 7.1 shows how much of the gap an optimized schedule recovers.

5. A conservative model in three dimensions

Why a thermal network rather than a flow solver? The evidence available to us constrains heat loss and body exchange, not velocity fields. A Navier–Stokes model would add parameters (a turbulence closure, the jet from the faucet) that nothing here can test, while hiding the quantity the decision depends on. We therefore keep the conservation laws exact and let a single effective diffusivity D stand for mixing, which keeps every assumption visible. The cost is that buoyancy and jet entrainment are not resolved.

The rectangular water envelope is divided into $8 \times 4 \times 3$ cells. Each has a capacity C_i and exchanges heat only across common faces. Top cells lose heat to the room; bottom and side faces lose heat through the shell. The bather is represented by a smooth, three-dimensional displacement field and a distributed skin contact term.

$$C_i \dot{T}_i = \sum_j g_{ij} (T_j - T_i) + H_{a,i} (T_a - T_i) + H_{b,i} (T_b - T_i) + S_i \quad (11)$$

$$g_{ij} = \rho c_p D \frac{A_{ij}}{d_{ij}} \min(\phi_i, \phi_j), \quad C_i = \rho c_p V_i \quad (12)$$

Here ϕ is fluid fraction, A_{ij} a face area and d_{ij} the center separation. Symmetric g_{ij} guarantees that internal heat exchange cancels when the equations are summed. D is an effective mixing closure; molecular diffusion alone would not represent motion or buoyant circulation.

$$w_i \propto \exp \left[-\frac{1}{2} \sum_{k=1}^3 \left(\frac{x_{ik} - b_k}{s_k} \right)^2 \right], \quad \sum_i w_i = 1 \quad (13)$$

$$V_i = V_{\text{cell}} - V_b w_i, \quad H_{b,i} = h_b A_b w_i \quad (14)$$

Baseline $b = (0.55L, W/2, 0.45H)$, with widths $s = (0.40, 0.16, 0.12)$ m. Fluid fraction is $\phi_i = V_i/V_{\text{cell}}$. The center b and widths s set the spatial distribution of body effects. Body volume and contact area remain independent inputs: shape is varied through s at fixed volume and area. This is a homogenized immersed-body representation, not an anatomically resolved obstruction. A cell whose occupied fraction reaches 0.95 is rejected rather than assigned a negative capacity.

A fixed envelope volume is preserved by balancing inlet and overflow. In the baseline the exact displaced-water volume is 164.25 L. Shape scenarios preserve envelope volume when isolating aspect-ratio effects. Volume scenarios deliberately change water depth and hence both storage and side-wall area.

5.1 Spatial evidence: the mean is not the whole bath




Figure 1. Final temperatures in the three horizontal cell layers. All layers use one color scale; the top layer contains the prescribed inlet-to-overflow stream. Geometry is in metres, and values are cell averages at $t = 30$ min from the Section 6 solver.

The heat map is calculated from the same archived trajectory as Figure 2. It shows where the imposed transport path and environmental/body sinks leave temperature differences. The plots are horizontal slices through a three-dimensional network with exchange between layers, not three independent two-dimensional models.

A temperature range summarizes the spread but cannot show its location. The maps make the physical interpretation intractable: localized replenishment and distributed losses must be balanced through mixing. Every cell contributes to the comfort test; a high mean cannot compensate for a cold region.

The Gaussian body representation changes both storage and exchange distribution. A map does not prove that this homogenized representation captures anatomy, recirculation or buoyancy. Its role is to reveal the actual implications of the declared model, not to create the appearance of a resolved flow simulation. The refinement table separately checks how cell size affects the result.

6. Inlet transport, solver and strategy search

Inlet and overflow are on opposite ends of a surface stream. Every path edge carries the same q : the first cell receives hot water, interior cells receive upstream water and lose the same volume downstream, and the last cell discharges through the overflow.

$$S_1 = \rho c_p q (T_{\text{in}} - T_1), \quad S_i = \rho c_p q (T_{i-1} - T_i) \quad (15)$$

$$\sum_i C_i \dot{T}_i = -\sum_i H_{a,i} (T_i - T_a) - \sum_i H_{b,i} (T_i - T_b) + \rho c_p q (T_{\text{in}} - T_{\text{out}}) \quad (16)$$

This stream deliberately allows hot-water short circuit. Heat can leave at a locally warm outlet before warming a distant cold region. It is an explicit transport assumption, not an inferred flow field. Alternative inlet placement would require another transport network and calibration.

For a fixed flow, the network is affine linear. Augmenting the state by a constant one permits matrix-exponential propagation, avoiding a large forward-Euler step restriction [4]. A delayed-start policy has two exact constant-control segments.

$$\dot{\mathbf{T}} = A(q)\mathbf{T} + \mathbf{b}(q), \quad \mathbf{z}(t + \tau) = e^{\tilde{A}(q)\tau} \mathbf{z}(t) \quad (17)$$

Starts are tested at 0, 120, ..., 1200 seconds. For each, rates from 0 to 3 L/min are bracketed on a 0.2 L/min grid; the first crossing of the lower-temperature requirement is refined by Brent's root method. The search targets $T_{\text{min}} + 0.03^\circ\text{C}$ as numerical reserve, then checks the upper-temperature and spread limits at 5-second output intervals. Starts that already violate the floor are rejected.

This is a reproducible candidate search. Nonmonotone flow responses, untested delays, pulsed inputs, inlet relocation and feedback are not excluded by the calculation. A failed search means no candidate was accepted, not that every possible action is infeasible. The independent checks further examine the chosen trajectory between samples.

7. The spatial result and the mixing penalty




Figure 2. Matrix-exponential cell temperatures under the selected policy. Cell temperatures are computed every 5 s by matrix exponentials; the shaded range spans the minimum and maximum over all cells, and the mean alone would hide cold and hot locations.

Result	Baseline
Rate / start	0.805 L/min / 0.0 min
Added water	24.14 L
Minimum / maximum	39.030 / 40.553°C
Maximum simultaneous spread	1.209°C
Final volume-weighted mean	39.418°C

Within the constant-rate family, starting the trickle immediately ranks ahead of delayed starts here, even though waiting is optimal in the ideal mixed model. The distant water needs time to receive heat; larger late rates create a warmer inlet region before they solve the cold-region constraint. This is a concrete consequence of resolving space.

Holding a perfectly mixed bath at 40°C throughout would use 36.22 L. This is a stricter reference service, not a like-for-like optimum. Our accepted 1°C cooling allowance uses less water partly because it provides a different service. The comparable mixed model with the

same lower limit uses 16.01 L. The gap also reflects our restricted control family and 0.03°C numerical reserve; it cannot be attributed purely to imperfect mixing.

7.1 A time-varying schedule uses less water

Constant-rate policies are a narrow family. We therefore optimize piecewise-constant inlet flow in K equal segments ($K = 3, 6, 12$) by sequential quadratic programming, minimizing added water under the same all-cell limits (with the 0.03°C reserve), sampled every 15 s and re-checked at 5 s, from four starting schedules. The problem is non-convex: each result is the best local optimum found, with no global claim.

Policy	Water (L)	Change	Min / max ($^{\circ}\text{C}$)	Spread ($^{\circ}\text{C}$)
Best constant rate	24.14	—	39.03 / 40.55	1.21
3 segments	19.98	-17.2%	39.03 / 40.84	1.50
6 segments	19.84	-17.8%	39.03 / 40.93	1.50
12 segments	19.77	-18.1%	39.03 / 40.98	1.50
Perfect-mixing optimum	16.01			
Energy lower bound	15.41			




Figure 3. Best 12-segment schedule against the best constant rate (top) and the range of cell temperatures (bottom). Flows come from multi-start sequential quadratic programming (Section 7.1) and temperatures are replayed every 5 s; dotted lines mark the limits.

The 12-segment schedule is off for the first 2.5 minutes and off for the last 10.0 minutes, with a peak of 1.57 L/min between. It adds 19.77 L, 18.1% below the constant rate, and narrows the gap to the energy lower bound from 8.7 L to 4.4 L. Refining from 3 to 12 segments gains only 1.1%: the saving comes from the shape. The optimum touches the lower limit and the spread limit (1.50°C) and comes close to the upper limit (40.98°C).

We have not isolated why the best schedule stops at the end. One candidate explanation, an inference rather than a model result, is that hot water added late leaves through the overflow before it warms the remote cells, while stored heat carries the final minutes.

8. Separate geometry, size and body effects

Change from baseline	Water (L)	Interpretation
Shallow Wide	No candidate	Same envelope volume
Deep Narrow	17.48	Same envelope volume
Small Bath	No candidate	Lower water depth
Large Bath	20.39	Higher water depth
Larger Body	No candidate	More displacement/contact
Long Body	23.85	Shape only; fixed volume/area
Warmer Skin	21.08	Skin at 36°C
Cooler Skin	27.18	Skin at 32°C

At equal volume, the wide, shallow tub has a larger exposed top area and a changed diffusion length, and no candidate was accepted for it. The deep, narrow tub needs 17.48 L against 24.14 L, because a deep, narrow envelope reduces top-area heat loss, although a real deep bath may develop vertical stratification not resolved by a constant D . Geometry is therefore more than an interchangeable cooling coefficient.

Changing depth at fixed length and width changes thermal storage and shell area. Greater water volume slows cooling, which can lower added water over a short fixed horizon: the large bath needs 20.39 L, while no candidate was accepted for the small bath. That does not mean a larger bath conserves total water: filling it initially uses more. J counts replenishment only, and this distinction prevents a misleading conservation claim.

The larger-body scenario simultaneously displaces more water and increases assumed contact area; it is a combined size scenario and has no accepted candidate. The long-body scenario isolates the spatial distribution by keeping total body volume and contact area fixed. Warmer skin reduces the temperature difference driving body heat loss. These changes quantify conditional dependence; they do not establish measured human heat transfer.

To transfer the model to a real tub, measure water volume after entry, submerged contact geometry, temperatures at several depths and distances, and no-inlet cooling. Shape and temperature should not be inferred from a single mean cooling curve.

9. Motion and a bubble layer can change the policy

Scenario	Added water (L)	Search outcome
Weak Mixing	—	None accepted
Strong Mixing	19.40	Accepted
Moving With Added Surface Loss	23.04	Accepted
Foam	10.37	Accepted
Tight Comfort	—	None accepted
Loose Comfort	1.87	Accepted
Low Loss	8.60	Accepted
High Loss	—	None accepted
Cool Supply	44.80	Accepted

Increasing D reduces the gradient created by the localized inlet. The strong-mixing scenario uses 19.40 L against 24.14 L, but motion may also increase heat loss. When D is increased together with a 20% increase in the surface coefficient (30 instead of 25 W/(m² K)), the need rises to 23.04 L and most of the benefit disappears. The comparison deliberately separates transport improvement from its possible boundary cost.

The foam scenario multiplies only the effective surface coefficient by 0.4. With all other inputs fixed, the best constant rate falls to 10.37 L. The assumed 60% reduction is a scenario, not an experimentally established property of bubble-bath additive. If the layer breaks up or motion raises evaporation, this result must be recomputed.

No accepted candidate exists for: Weak Mixing, Tight Comfort, Shallow Wide, Small Bath, Larger Body, High Loss. They are reported as such, not hidden in a favorable average. Their outcomes identify where the restricted strategy must change. Potential responses include improved circulation, a different inlet path, a shorter bath, or a different comfort tolerance. The paper does not certify which response is safe or optimal without corresponding physical evidence.

10. Validation: independent evidence and remaining error




Figure 4. Water-use comparison under the same lower limit. The energy bound is conditional on the network assumptions; the feasible spatial result remains above it.

All 17 recorded checks pass. Independent heat-flow arithmetic sums environmental losses, body exchange and local-temperature overflow; the largest instantaneous residual is $2.46\text{e-}10$ W. An adaptive RK45 integration with independently assembled right-hand side agrees with the archived matrix-exponential trajectory within $3.12\text{e-}07^\circ\text{C}$. This tests numerical implementation, not a second physical model.

A constructed uniform-loss problem reproduces its analytic exponential cooling solution. Internal exchange cancels globally, and the stream has exactly one overflow sink. Positive off-diagonal dynamics and nonpositive row sums support physical bounds and a contractive derivative estimate for each constant-flow segment.

One-second sampling is supplemented by a derivative bound: if $|dT_i/dt| \leq L$, a point lies at most 0.5 seconds from a sample, so temperatures differ by at most $L/2$ and spread by at most L . With a numerical allowance, the resulting lower envelope is 38.9989°C , upper 40.5839°C , and spread upper bound 1.2715°C . This is a conditional floating-point envelope, not interval-arithmetic certification.

The optimized schedule is replayed with the independently assembled right-hand side and RK45 at one-second samples. It agrees with the matrix-exponential trajectory within $3.3\text{e-}07^\circ\text{C}$ and satisfies the limits (minimum 39.030°C , maximum 40.984°C , spread 1.500°C , each within the 0.002°C sampling tolerance); its 19.77 L exceeds the 15.41 L energy bound, as it must. These two checks cover this one schedule and are separate from the 17 model checks.

11. Resolution, uncertainty and transfer to a real bath

Resolution / replay	Water (L)	Minimum	Maximum spread
$8 \times 4 \times 3$ selected	24.142	39.030°C	1.209°C
$12 \times 6 \times 4$ selected	24.117	39.030°C	1.416°C
Coarse policy on fine grid	24.142	39.031°C	1.417°C

The selected water amounts differ by 0.11%. The fine grid preserves the coarse-policy constraints at sampled times. The larger fine-grid spread exposes a more localized inlet temperature: a close water-volume result does not establish pointwise temperature convergence. Two meshes are a diagnostic, not an asymptotic convergence-order study.

Parameter uncertainty concerns heat loss, body exchange and mixing. Structural uncertainty concerns the prescribed surface stream, fixed skin reservoir, homogenized body and constant effective D. These are different errors. The scenario table explores parameter dependence and one coupled motion effect; it is not a probability distribution or confidence interval.

Literature-range analysis. Section 2 gives ranges for six inputs. 64 scrambled Sobol points cover them (D log-uniform), and the constant-rate search runs at each. 37 points have an accepted candidate. Among those, the required water is 12.5, 16.5, 20.8, 26.9 and 31.1 L at the 5th, 25th, 50th, 75th and 95th percentiles. The remaining 27 points, mostly at high loss, have none. The baseline of 24.1 L sits in the upper half of this distribution.

Spearman rank correlations with water, over accepted draws: surface coefficient +0.65, room temperature -0.49, body coefficient +0.46, skin temperature -0.29, shell coefficient +0.16, mixing diffusivity -0.09.

The correlations use accepted draws only, so they understate inputs that push draws into infeasibility. The draws are uniform over ranges that we chose; they are not a probability distribution for any real tub. The main message is that the loss coefficients and the room temperature, not the mixing closure, decide how much water is needed, which is why Section 2 anchors them in outside evidence.

A calibration plan should begin with an inlet-free cooling experiment, recording water volume and room conditions. Multiple temperature probes distinguish aggregate heat loss from mixing. A separate inlet experiment records flow and inlet/outlet temperatures, allowing overflow energy to be checked. Motion and foam need matched trials because both can alter transfer coefficients. Reserve an entire experiment for prediction checks after fitting. A single mean-temperature series cannot separate the surface, shell and body coefficients, so only calibrated and independently tested parameters would support advice for a particular bath.

12. Conclusions and a policy with clear scope

The model supports four linked findings. First, tolerating a modest temperature decline lets stored heat replace some inlet water. Second, spatial transport changes the ideal mixed recommendation: the best constant rate starts at once, and a schedule that withholds water at both ends needs 18% less. Third, the water required is dominated by the surface and body loss coefficients: from 12 to 31 L across the literature-based ranges. Fourth, improving mixing helps only to the extent that it does not create offsetting boundary loss or hot-water short circuit.

For the stated baseline, the best constant trickle is 0.805 L/min for 24.14 L over 30 minutes; the 12-segment schedule uses 19.77 L. These are the best accepted results of the documented searches, with all-cell constraints and independent numerical checks. Neither is a universal faucet prescription or a proved global optimum over all time-varying controls.

The spatial model earns its complexity by explaining a disagreement with a proved simple benchmark and by rejecting strategies that a mean-only calculation would accept. Conservative accounting, anchored coefficients with stated ranges, a feasible-policy versus lower-bound gap, and an explicit refinement check make the recommendation inspectable.

The limitations are consequential. Three-dimensional geometry is represented, but fluid momentum, buoyancy, jet entrainment, free-surface motion and detailed body anatomy are not solved. No bath experiment was conducted. Temperature limits express a preference assumption. Coarse cell averages cannot establish burn safety near the inlet, and the shape scenarios do not resolve stable vertical layers.

Before using a numerical rate in practice, identify the real tub's cooling and mixing behavior. Until then, the defensible transferable advice is to avoid unnecessary replenishment, distinguish cold-region temperature from the mean, improve distribution before increasing flow, and reconsider a strategy when geometry or motion changes. The following page translates these principles without requiring the user to interpret the equations.

13. A warmer bath, with less replacement water

A guide for the person in the bathtub

Decide what “warm enough” means. A bath need not stay at exactly its starting temperature to remain acceptable. In our example, allowing a 1°C fall still needs about 24 L of replacement water, while allowing a 2°C fall needs only 1.9 L. Your actual comfort and health needs must determine the acceptable range.

Distribute the warmth before turning up the tap. The water near the faucet can warm while the far end remains cool. Gentle movement can help redistribute heat. More vigorous motion is not automatically better: it may also increase heat loss from the surface.

Run the tap in the middle, not at the ends. For the tub and mixing conditions analyzed here, the least-water schedule adds no hot water for the first 2.5 minutes or the last 10 minutes and a steady stream in between; a constant trickle from the start used 22% more. In a bath that mixes very quickly, waiting until it has cooled slightly can save more. There is no single schedule that works for every tub.

Check more than one location. A comfortable mean temperature can hide a cold far corner or a hot inlet region. Check the bath away from the faucet and at different depths. Do not use a model’s cell-average temperature as a reason to contact a hot jet.

Remember where the added water goes. Once the tub is full, extra water leaves through the overflow. Hot water that runs along the surface to the overflow may escape before its warmth reaches you. Increasing the stream can waste heat without fixing unevenness.

A bubble layer may help, but the amount is uncertain. An intact insulating surface layer can reduce heat loss. Our calculation explores such a layer; it does not measure the effect of a particular additive. If the layer disappears or the bath behaves differently, reassess rather than relying on the example’s saving.

Ideas the model does not test. Pouring the hot water in at the end of the tub farthest from the drain, draining a little cooler water before adding hot water, or stirring from the bottom could each help, because the model shows that the coldest region limits the plan. Each needs a different transport network and evidence, so none is claimed here.

The practical rule: use the bath’s stored warmth, keep temperature distributed, and add only as much water as your actual conditions require. The numerical example explains these tradeoffs; it is not a tested bathing or safety standard.

14. References and reproducible algorithm

- [1] COMAP. 2016 MCM Problem A: A Hot Bath. 2016. Official problem PDF: contest.comap.com/undergraduate/contests/mcm/contests/2016/problems/2016_MCM_Problem_A.pdf.
- [2] U.S. Department of Energy. DOE Fundamentals Handbook: Thermodynamics, Heat Transfer, and Fluid Flow. DOE-HDBK-1012/2-92, June 1992. Convection heat transfer, Eq. (2-9).
- [3] IAPWS. Revised Release on the IAPWS Formulation 1995 for the Thermodynamic Properties of Ordinary Water Substance for General and Scientific Use. R6-95(2018). iapws.org/technical-guidance/release/IAPWS-95.
- [4] SciPy developers. `scipy.sparse.linalg.expm_multiply` and `scipy.integrate.solve_ivp`, API documentation. docs.scipy.org/doc/scipy/reference/. Accessed 7 October 2026.
- [5] Shannon, K. M. Judges' Commentary: Hot Bath Problem. The UMAP Journal 37(3), 277–281, 2016. faculty.winthrop.edu/abernathy/MathComps/2016_MCM_A-Com.pdf.
- [6] Bergman, T. L., Lavine, A. S., Incropera, F. P., DeWitt, D. P. Fundamentals of Heat and Mass Transfer, 7th ed. Wiley, 2011. Natural convection above a heated horizontal surface (characteristic length A/P); heat and mass transfer analogy.
- [7] Menzies, C., Clarke, N., Steward, C. J., Thake, C. D., Pugh, C. J. A., Cullen, T. Vascular, inflammatory and perceptual responses to hot water immersion: impacts of water depth and temperature in young healthy adults. *Experimental Physiology*, 2025. doi:10.1113/EP092761.
- [8] OpenAI, Codex (GPT-6-based assistant), and Anthropic, Claude Sonnet 5.5 in Claude Code. Used on 7 October 2026 for modeling, code, validation and report composition; exact builds not independently established. See Report on Use of AI Tools.

Algorithm and computational record

1. Form cell volumes and capacities; reject excessive body occupancy. Assemble symmetric face exchange, environmental/body conductances and the conservative stream. 2. Compute the mixed closed forms and energy lower bound. 3. Evaluate delayed constant-flow candidates with matrix exponentials and refine the first bracketed lower-temperature crossing. Reject spread or upper-limit violations. 4. Rank accepted candidates by added litres. 5. Independently integrate and check balances, analytic limits, parameter arithmetic and constraints; replay a finer mesh and perturb physical scenarios. 6. Optimize piecewise-constant schedules (multi-start SQP) and replay the best independently. 7. Rerun the constant-rate search at Sobol draws over the literature ranges. Independent RK45 checks use relative tolerance $2e-9$, absolute $2e-10^{\circ}\text{C}$ and steps of at most 5 s. The organization and practical page respond to the guidance in [5]; AI assistance is acknowledged in [8].

Appendix A. Scenario definitions

Each row changes only the listed baseline inputs. All unspecified inputs remain as stated in Section 2; geometry and conductance are then rebuilt. The equal-volume aspect-ratio cases retain the envelope volume of 0.22425 m^3 . Widths s describe a smooth displacement distribution, not measured anatomical semiaxes.

Scenario	Changed input(s)
Weak Mixing	$D \text{ (m}^2/\text{s)} = 0.0003$
Strong Mixing	$D \text{ (m}^2/\text{s)} = 0.003$
Moving With Added Surface Loss	$D \text{ (m}^2/\text{s)} = 0.003$; surface coefficient ($\text{W}/\text{m}^2\text{K}$) = 30
Loose Comfort	lower limit ($^{\circ}\text{C}$) = 38
Tight Comfort	lower limit ($^{\circ}\text{C}$) = 39.5
Foam	surface multiplier = 0.4
Shallow Wide	length (m) = 1.7; width (m) = 0.75; depth (m) = 0.175882
Deep Narrow	length (m) = 1.3; width (m) = 0.6; depth (m) = 0.2875
Small Bath	depth (m) = 0.2

Accepted candidates are independently replayed on their scenario network at intervals of no more than five seconds. These scenario checks are sampled checks; the stronger one-second continuous-time envelope is reported for the selected baseline policy only. “None accepted” refers to the documented finite search and is not an infeasibility theorem.

Appendix A. Scenario definitions (continued)

Each row changes only the listed baseline inputs. All unspecified inputs remain as stated in Section 2; geometry and conductance are then rebuilt. The equal-volume aspect-ratio cases retain the envelope volume of 0.22425 m^3 . Widths s describe a smooth displacement distribution, not measured anatomical semiaxes.

Scenario	Changed input(s)
Large Bath	depth (m) = 0.27
Larger Body	body displacement (m^3) = 0.075; body contact area (m^2) = 1.35
Long Body	body widths s (m) = [0.48, 0.12, 0.1]
Warmer Skin	skin temperature ($^{\circ}\text{C}$) = 36
Cooler Skin	skin temperature ($^{\circ}\text{C}$) = 32
High Loss	surface coefficient ($\text{W}/\text{m}^2\text{K}$) = 36; wall coefficient ($\text{W}/\text{m}^2\text{K}$) = 8.5; body coefficient ($\text{W}/\text{m}^2\text{K}$) = 40
Low Loss	surface coefficient ($\text{W}/\text{m}^2\text{K}$) = 17; wall coefficient ($\text{W}/\text{m}^2\text{K}$) = 4.5; body coefficient ($\text{W}/\text{m}^2\text{K}$) = 12
Cool Supply	inlet temperature ($^{\circ}\text{C}$) = 45

Accepted candidates are independently replayed on their scenario network at intervals of no more than five seconds. These scenario checks are sampled checks; the stronger one-second continuous-time envelope is reported for the selected baseline policy only. “None accepted” refers to the documented finite search and is not an infeasibility theorem.

Report on Use of AI Tools

Tool and scope

OpenAI Codex, a GPT-6-based assistant, was used on 7 October 2026. The later revision (literature-based coefficients, schedule optimization, range analysis and rebuilding of this report) used Anthropic's Claude, Sonnet 5.5, through Claude Code on the same date. The assistant selected the historical problem, located and read sources, formulated the mixed and spatial models, wrote and revised Python code, designed checks, interpreted numerical runs, generated figures, and drafted and typeset the English report. An assistant sub-agent wrote the independent-RHS validation routine; this is AI-assisted code review, not review by another human.

Task

The task was to select a classic historical MCM problem and build a rigorous, professional case from it. The conversation wording is not reproduced in this report or in the development record.

Outputs, corrections and verification

AI-produced outputs are the model, validation and report-building sources; the mathematical derivations, scenario tables, figures and text in this document. Numerical values came from actual Python runs. A weak-mixing policy search failed the imposed constraints and remains reported as such. An early continuous-time envelope failed; the flow-search reserve was increased from 0.01°C to 0.03°C before the accepted run. Citations [1]–[5] were checked against the linked sources; [6] was not opened (its correlations were checked against web summaries) and [7] was read as an abstract only. The first baseline coefficients and the 40.5°C upper limit were replaced after anchoring and mesh refinement.

Record limitations

A full exported interaction transcript and every intermediate AI output were not available in this artifact. The task description, tool identification, scope, code artifacts and numerical receipts are retained in the accompanying development record; this report does not invent a transcript. No independent human review or physical bath experiment is claimed.

Python, NumPy and SciPy performed the mathematics; Matplotlib generated plots and STIX vector formula outlines; ReportLab produced the PDF with embedded Times New Roman faces. These deterministic tools do not establish empirical validity. The service model/build beyond the available GPT-6-based identifier was not independently verified.